

A Simple MNIST Digit Classifier Neural Network Made from scratch in Python

Thiruwaran Kalvin (T-Kalv)
Undergraduate Computer Science Student

August 18, 2025

1 Abstract

This research paper presents an implementation of a basic feed-forward neural network for classifying/recognising handwritten digits from the MNIST dataset, built from scratch in Python using only core libraries such as NumPy, Matplotlib, Pandas, and tqdm. This feed-forward neural network model, also known as a multilayer perceptron (MLP), achieves a classification probability accuracy of approximately 89% for both the training and testing datasets. This demonstrates the viability of simple neural networks that recognise handwritten numerical digits without the need for external heavy libraries such as PyTorch or TensorFlow, instead utilising only core Python libraries and Jupyter Notebook.

2 Introduction



Figure 1: Automatic processing of handwritten bank cheque images: A survey. International Journal on Document Analysis and Recognition

Numerical Digit Classification is a specific field of Image Processing that focuses here on this paper, classifying and recognising handwritten low-level black and white numerical digit images using the MNIST Dataset[1]. This particular task of recognising handwritten digits is vital in the real-world as previously Bank Tellers played a crucial role in processing Cheques by manually verifying (looking) at the authenticity of the Cheques and verifying the Currency amount on the Cheques (as illustrated in 'Figure 1 - Automatic processing of handwritten bank cheque images: A survey. International Journal on Document Analysis and Recognition'). This method was time-consuming and in-efficient due to manual Human validation as well as being prone to Human errors. As the level of technology/computation advanced, the need to accurately read hand-written numerical digits especially in sectors such as the Banking industry became more crucial to reduce the likelihood of errors and reduce processing time. This results in Financial institutions diverting resources more effectively to other services. Through the advancements in Artificial Intelligence and Machine learning, we can develop systems that can accurately recognise hand-written digits much more efficient and effectively using hand-written datasets such as MNIST.

MNSIT (Modified Natual Institute of Standard and Technology) database is dataset that consits of handwritten numerical digits that is commonly used for image processing training and contains 60,000 training images used to train the Machine Learning Neural Network model and 10,000 testing images to validate the accuracy of the model. In this paper, we use this famous MNIST dataset which is the standard benchbark for Image Classification algorithms. As the MNIST database was constructed before the turn of the 21st Century, there have many implementations/existing solutions using the MNIST Dataset to classify and recognise hand-written numerical digits such as using tradition methods (including but not limited to KNN (K-Nearest Neighbours), SVM (Support Vector Machines)) and even Deep-Learning methods (including but not limited to CNN (Convuntional Neural Network), RNN (Recurrent Neural Network)). Arguably, one of the most famous implementations is the LeNet-5 architecture [3] which uses CNN consiting of 2 convolutions, 2 subsamplings and 3 fully connected layers with around 60000 trainable paramaters.

This paper proposes a light-weight simple MNIST Digit Classifier Neural Network [4] consiting of 2 layers with 784 neurons, built without the use of external heavy deep-learning libraries (only using libraries such as NumPy, Matplotlib) with a particular emphasis on understanding and manually implementing algorithms and functions. This work presents such algorithms and functions such as Forward Propagation, Backpropagation, ReLu (Rectified Linear unit) activation function, Softmax activation function.... This results in a reproducible implementation of simple MNSIT Digit Classifier Neural Network, with clear, hand-on demonstration of 'behind-the-scenes' of the network as well as the training procedure and performance benefits/outcomes.

This paper is organised as follows: Section 3 discusses related work, followed by Section 4 which outlines the Methodology and architecture, Section 5 presents Experimental Evaluation and finally Section 6 concludes with findings and future work.

3 Related Work

Before working on this 'Simple MNIST Digit Classifier Neural Network', previously a 'Simple Neural Network' was developed from scratch in Python for XOR [4]. This previous project implemented key concepts such as Forward Propagation, Backpropagation, Mean-Squared Error Loss and the Sigmoid Activation function using only core libraries such as NumPy and Matplotlib. The sucess of XOR having a probability of approximately 95% helped to inspire the next challenge of handling

a more real-world dataset MNIST, implementing these key concepts on this 'Simple MNSIT Digit Classifier Neural Network'.

Previous research into using the MNIST dataset for classifying and recognising handwritten digits has produced a wide-range of models with a variety of complexity and performance benefits to the traditional manual, human-validation of handwritten digits. LeNet-5 [3], introduced in 1998 having achieved approximately a 1%[5] test error probability which resulted in approximately 99% probability accuracy with the MNSIT dataset, implemented using CNN (Convolutional Neural Networks). While LeNET-5[5], achieved a respectable, high probability accuracy using CNN, Max Pooling with 2 convolutions, 2 subsampling, 3 fully connected layers and around 60000 trainable paramaters, this paper investigates the performance of a simple multilayer perceptron (MLP) using a feed-forward neural network trained from scratch without the need of heavy, high-level frameworks. However, more modern implementations do use these heavy, high-level frameworks such as PyTorch [6], Keras [7], TensorFlow [8] utilising the performance benefits of GPUs (Graphics Processing Unit), NPU (Neural Processing Unit), TPUs (Tensor Processing Units) which help to optimise the speed of learning/training as well with the use of advanced optimisers (including but not limited to Adam). Here due to the added cost of utilising such performance enhancers, the system here presented in this paper will instead utilise modern CPUs for computation and training.

4 Methodology

4.1 Dataset




Figure 2: Sample handwritten digits from the MNIST dataset. Image by Suvanjanprasai, CC BY-SA 4.0.

The MNIST dataset [9] contains gray-scale, black and white images of hand-drawn numerical digits from zero to nine (0, 1, 2, 3, 4, 5, 6, 7, 8, 9). Each image is 28 pixels in height by 28 pixels in width) contains a total of 784 pixels, where each pixel value is an integer from 0 to 255 resulting in black to white. The dataset consists of 60000 training images for training the model and 10000 testing images for testing the trained model. The training dataset in this paper has 785 columns which are the 'label'. This is the corresponding '0 to 9' digit that was handwritten and the remaining columns are the '0 to 255' pixel values associated with the black and white image. The testing dataset is the same as the training dataset except without containing the labels to validate the accuracy of the model.

4.2 Network Architecture




Figure 3: Simplified Neural Network Architecture used in this Project

The Neural Network [Figure 2] is a fully-connected, feed-forward neural network that consists of an input layer, hidden layer and output layer. The input layer (first layer x) consists of 784 neurons which corresponds to the 784 pixels in the 28×28 flattened image for each of the images in the MNIST dataset. The hidden layer (second layer) consists of 10 neurons where the corresponding weights and biases connecting to this layer and initialised with random values between -0.5 to 0.5. In addition, the activation function used in the hidden layer is the ReLU (Rectified Linear unit) activation function. The output layer (third layer) consists of 10 neurons which is the prediction of what numerical digit the input image represents from 0 to 9 resulting in 10 possible numerical digit classes.

The parameters of this Neural Network consists of a 'weight1' which has a shape of (10,784) connecting the input to the hidden layer, 'bias 1' which has a shape of (10,1) which are the biases for the hidden layer, 'weight 2' which has a shape (10,10) connecting the hidden layer to the output layer and 'bias2' which has a shape of (10,1) which are the biases of the output layer. This results in this simple MNIST Digit Classifier Neural Network having a total of 7960 trainable parameters with 7840 weights, 10 biases between the input layer and hidden layer, 100 weights, 10 biases between the hidden layer and output layer.

4.3 Forward Propagation

1: Forward Propagation Pseudocode Algorithm

```
FUNCTION forwardPropagation(weight1, bias1, weight2, bias2, X):  
    Z1 = matrixMultiply(weight1, X) + bias1  
    A1 = ReLuFunc(Z1)  
    Z2 = matrixMultiply(weight2, A1) + bias2  
    A2 = softMaxFunc(Z2)  
    RETURN Z1, A1, Z2, A2
```

2: ReLu Activation Function Pseudocode

```
FUNCTION ReLuFunc(Z):  
    RETURN MAXIMUM(0, Z)
```

3: Softmax Activation Function Pseudocode

```
FUNCTION softmaxFunc(Z):  
    maxZ = MAX(Z, axis=0)  
    shiftedZ = Z - maxZ  
    expZ = EXP(shiftedZ)  
    sumExpZ = SUM(expZ, axis=0)  
    RETURN (expZ / sumExpZ)
```

The Forward Propagation algorithm used in this project, takes the input data image and runs through the Neural Network layers to determine what the output is going to be. Here in this 'Simple MNIST Digit Classifier Neural Network' the input data is a flattened 28x28 pixel image consisting of 784 pixels. This input image is then represented as a matrix which each row is an example with 784 columns corresponding to 1 pixel in the image. This matrix is then transposed resulting in the rows examples becoming columns examples.

In the first layer (unactivated layer), the input 'X' is matrix multiplied by the weight1 (first weight) and then the bias term is added, resulting in Z1 which is the first unactivated first layer in the matrix (as illustrated in '1: Forward Propagation Pseudocode Algorithm' above).

The second layer is a linear combination of the first layer which results in a purely linear transformation, hindering the Neural Network's ability to model non-linear relationships. This leads to an application of a ReLu (Rectified Linear unit) activation function which is applied to Z1 to introduce non-linearities where it processes each element in Z returning the value if it greater than zero or return zero otherwise (as illustrated in '2: ReLU Activation Function Pseudocode' above) leading to A1.

In the pre-activation, output layer, A1 is then matrix multiplied by weight2 (second weight) and then the second bias term is added, resulting in Z2. This leads to an application of the Softmax activation function (as illustrated in '3: Softmax Activation Function Pseudocode' above) which is used to convert the raw values into probability values for each of the 10 possible numerical digit classes (0, 1, 2, 3, 4, 5, 6, 7, 8, 9), resulting in A2.

4.4 Backpropagation

The Backpropagation algorithm used in this project, uses the predicted output probabilities from the Forward Propagation algorithm as a starting point. The Backpropagation algorithm then computes the gradients of the loss function with respect to each of the parameters in the Neural Network. It does this by calculating how much the Forward Propagation prediction deviated by and compares it to the actual label from the MNIST dataset resulting in an error loss value and calculates how much the previous weights and biases contributed and adjust these values, from the output layer moving backward through the network in the opposite direction to the movement of Forward Propagation algorithm.

$$\text{one_hot}(i, K) = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \\ \vdots \\ 0 \end{bmatrix}, \quad \text{where the 1 appears at index } i \text{ and } K \text{ is the total number of classes.}$$

Figure 4: One-Hot Encoding Algorithm

Firstly, a One-Hot Encoding algorithm is applied (as illustrated in 'Figure 3: One-Hot Encoding Algorithm') where it creates a matrix array of zeros with 10 output classes, 0 to 9, and goes through each row in the matrix specified by the label Y and sets it to 1. Then it tranposes the matrix so that each coloumn is an example instead of a row.

After the One-Hot Encoding Matrix Y is returned, the output layer calculation is performed. This is done by first subtracting the One-Hot Encoding Matrix Y labels from the predicted probabilities from the Forward Propagation algorithm. This reults in dZ2 which is error of the second layer in the Nerual Network, which represents the deviation from the correct label.

dW2, which is the weights gradient in the second layer, is calculated by first calculating the size of the One-Hot Encoding Matrix Y, m. Then it multiplies dZ2 with the transpose of A1 (which is the hidden layer activations) and divided by m. This results in the derivative of the loss function with respect to the weights in layer 2.

dB2, which is the bias gradient for the second layer in the Neural Network, is computed by getting the size of the One-Hot Encoding Matrix Y, m and then then summing all of the elements in dZ2 and dividing by m. This results in average of the absolute error for the second layer, dB2.

After the average of the absolute error for the second layer is obtained, the error is then back propagated to the hidden layer by multiplying the transpose of weight2 with dZ2 and then multiplying by the ReLu derivative activation function to Z1. This results in dZ1 which represents how much the first hidden layer has been deviated by.

Then DW1, which is the first layer weight gradient, is computed by getting the size of the One-Hot Encoding Matrix Y, m and then multiplying dZ1 with the input Matrix X that has been

transposed and divide by m . This results in the how much weight1 in the first layer contributed to next layer's (hidden layer) error.

Finally, $dB1$, which is the first layer biases gradient, is computed by getting the size of the One-Hot Encoding Matrix Y , m and then summing all the elements in $dZ1$ and dividing by m . This results in how much bias1 contributed to the first layer to the Neural Network overall error.

4.5 Training

4: Initialise Paramaters Pseudocode

```
FUNCTION initialiseParameters():
    weight1 = randomMatrix(10, 784) - 0.5
    weight2 = randomMatrix(10, 10) - 0.5
    bias1 = randomMatrix(10, 1) - 0.5
    bias2 = randomMatrix(10, 1) - 0.5
    RETURN weight1, weight2, bias1, bias2
```

5: Update Paramaters Pseudocode

```
FUNCTION updateParameters():
    weight1 = weight1 -  $\alpha$  x dW1
    weight2 = weight2 -  $\alpha$  x dW2
    bias1 = bias1 -  $\alpha$  x dB1
    bias2 = bias2 -  $\alpha$  x dB2
    RETURN weight1, weight2, bias1, bias2
```

In order for this simple MNIST Digit Classifier Neural Network to learn from the training dataset, a gradientDescent function is used where it iteratively adjusts the Neural Networks parameters (weights and biases) to accurately predict numerical handwritten digits based on errors made before.

Firstly, we assign random starting values to weight1, weight2, bias1, bias2 in each of the layers and subtract by 0.5 which helps to ensure the initial weights and biases are centered around zero than being positive (as illustrated in '4: Initialise Paramaters Pseudocode' above). This results in the Neurons in the Neural Network start with different parameters while training and ensures the Network learns effectively.

After the paramters have been initialised, forwardPropagation is performed where it takes the input data image and runs through the Neural Network layers to determine what the output data is going to be (as mentioned earlier in this paper in 'Section 4.3 - Forward Propagation').

The predicted ouput probabilities from forwardPropagation is then used in Backpropagation, where it computes the gradients of the loss function with respect to each of the parameters in the Neural Network and calulated the deviation of its predictions from the correct label from the MNIST dataset (as mentioned earlier in this paper in 'Section 4.4 - Backpropagation').

After the gradients of the loss function with respect to each of the parameters in the Neural Network have been computed, updateParameters is performed where each weight1, weight2, bias1,

bias2 are adjusted depending on the results in the gradientDescent algorithm. Here in updateParameters, each parameter is updated by getting the product of the rate of learning α and subtracting it with its corresponding gradient. This process is repeated for each of the parameters where it gradually reduces the error loss by moving the parameters in the direction where it improved the performance of the Neural Network.

Finally, here in this simple MNIST Digit Classifier Neural Network Model, for every 10th iteration we fetchPrecisions and fetchAccuracy while training.

5 Experimental Evaluation

5.1 Benchmarks and Setup

In order to evaluate the performance of the 'Simple MNIST Digit Classifier Neural Network', the MNIST dataset for classifying and recognising handwritten digits was used as the benchmark. This MNIST dataset consists of 60000 training images for training the model and 10000 testing images for testing the trained model. Each image is 28x28 black and white pixel image with a total of 784 pixels per image, labelled with the corresponding digit classifying '0 to 9' digit. This benchmark was selected for this Research Project as it serves as the main benchmark for evaluating image processing systems/ machine learning algorithms such as for this 'Simple MNIST Digit Classifier Neural Network' as well as allowing for direct comparison with the LeNet-5 architecture (LeCun et al., 1998) [3].

All experiments were conducted on a Laptop equipped with Intel i7-1360P CPU (5GHz), 16 GB RAM, running Ubuntu 25.04. The implementation was programmed using Python 3.13.3 via Jupyter Notebook using the core libraries: Numpy 2.3.2, Matplotlib 3.10.5, Pandas 2.3.1, tqdm 4.67.1. All training was performed on-device via CPU without GPU Or TPU Or NPU acceleration.

5.2 Results and Observations

The first 'working implementation*' achieved a classification probability accuracy of approximately 9.85% on the training dataset with 100 iterations and rate of learning of 0.1. This resulted in accuracy roughly corresponding to randomly guessing which was not the intention of this experiment. After post-investigation, multiple implementation were revealed to be the cause of the low classification probability accuracy. The main issue here was an incorrect Softmax activation function implementation where initially the function sums the raw score Z and not the exponentials as well as summing all the entries at once instead of each sample one normalisation. This was corrected by normalising across for each sample as well as subtracting by the maximum per sample for stability.

The second 'working implementation*' achieved a slightly improved classification probability accuracy of approximately 11% on the training dataset with 500 iterations and rate of learning of 0.1. After post-investigation, multiple implementation issues were discovered. The first main issue here was related to implementing the wrong update gradient where initially 'weight2 = weight2 - α * dB2' had been implemented instead of multiplying by dW2 for the update weight in the second layer which resulted in slow rate of learning. The second main issue here was that the inputs were normalised instead of feeding raw pixel values from 0 to 255 which also resulted in the training of the Neural Network slow and hurt convergences. This was resolved by scaling the inputs to $[0, 1]$, 'trainImages = trainImages / 255, devImages = devImages / 255'.

The third 'working implementation*' achieved a much improved classification probability accuracy of approximately 89% on the training dataset with 1000 iterations and rate of learning of 0.1. In the testing dataset, it also achieved a much improved classification accuracy of approximately 88% with 1000 iterations and rate of learning 0.1. This much improved implementation was due to the fixes previously mentioned before as well as implementing a progress indication (via the tqdm library) during training so that the user can track the progress throughout the training of the Neural Network model. These changes resulted in a more stable learning experience for the model, correctly identifying 10 random numerical handwritten digits out of the possible 10 random numerical handwritten digits from the training dataset when testing after training the model with the training dataset. With the testing dataset, out of possible 10 random numerical handwritten digits from the testing dataset, the trained Neural Network model achieved correctly 9 numerical handwritten digit classifications.

5.3 Limitations and Anomalies

While repeated training of the 'Simple MNIST Digit Classifier Neural Network' and evaluating results, random, non-deterministic misclassification results were observed. For example, during one run of training at the end [10], the Neural Network predicted the numerical digit '3' instead of correctly predicting the numerical digit '8', resulting in a misclassification. We cannot fully explain this behaviour but we think it is due to the randomness in the training process such as initialising random weights, shuffling the training set, the precision of floating point in Matrix operations. These stochastic effects may have led to slight variations in parameters while learning, resulting in this misclassification.

6 Conclusion and Future Work

6.1 Summary

We have presented the 'Simple MNIST Digit Classifier Neural Network', a system that is made from scratch in Python with only core libraries such as NumPy, Matplotlib, Pandas and tqdm. The purpose was to implement a feed-forward neural network model/MLP (multilayer perceptron) that can recognise and classify numerical handwritten digits from the MNIST data as well as gaining a deeper understanding of the underlying core mathematics and algorithms such as forward propagation, backpropagation, ReLu (Rectified Linear unit) activation function, Softmax activation function...

6.2 Experimental Results

The experimental results obtained using the MNIST dataset indicate that the 'Simple MNIST Digit Classifier Neural Network' achieved an accuracy of 89% with consistent performance throughout the training of the network and testing of the network despite occasional stochastic results. All experiments were conducted with only CPU computation without access to a GPU, NPU, TPU hardware acceleration as well as no use of heavy, high-level frameworks such as PyTorch, TensorFlow, JAX. This clearly demonstrates the viability that a model can achieve respectable performance on-par with modern implementations and previous implementations with only core libraries and fundamental mathematics.

6.3 Conclusion

In conclusion, the results clearly demonstrate the viability that a model can achieve respectable performance on-par with modern implementations and previous implementations with only core libraries and fundamental mathematics without leveraging performance enhancers such as high-level machine learning frameworks and hardware acceleration. However, the occasional stochastic misclassification of numerical handwritten digits that were observed indicate that model stability is an important consideration in the future. Furthermore, the current design of this simple neural network hinders its ability to capture more complex classification is also an important consideration for the future.

6.4 Future Work

Future work would be to further investigate the non-deterministic misclassifications by temporarily fixing random seeds to increase the reproducibility of results and increasing the overall performance of the model to achieve 99% coverage across the MNSIT dataset. This can be done by implementing convolutional layers, regularisation techniques...

Appendix

References